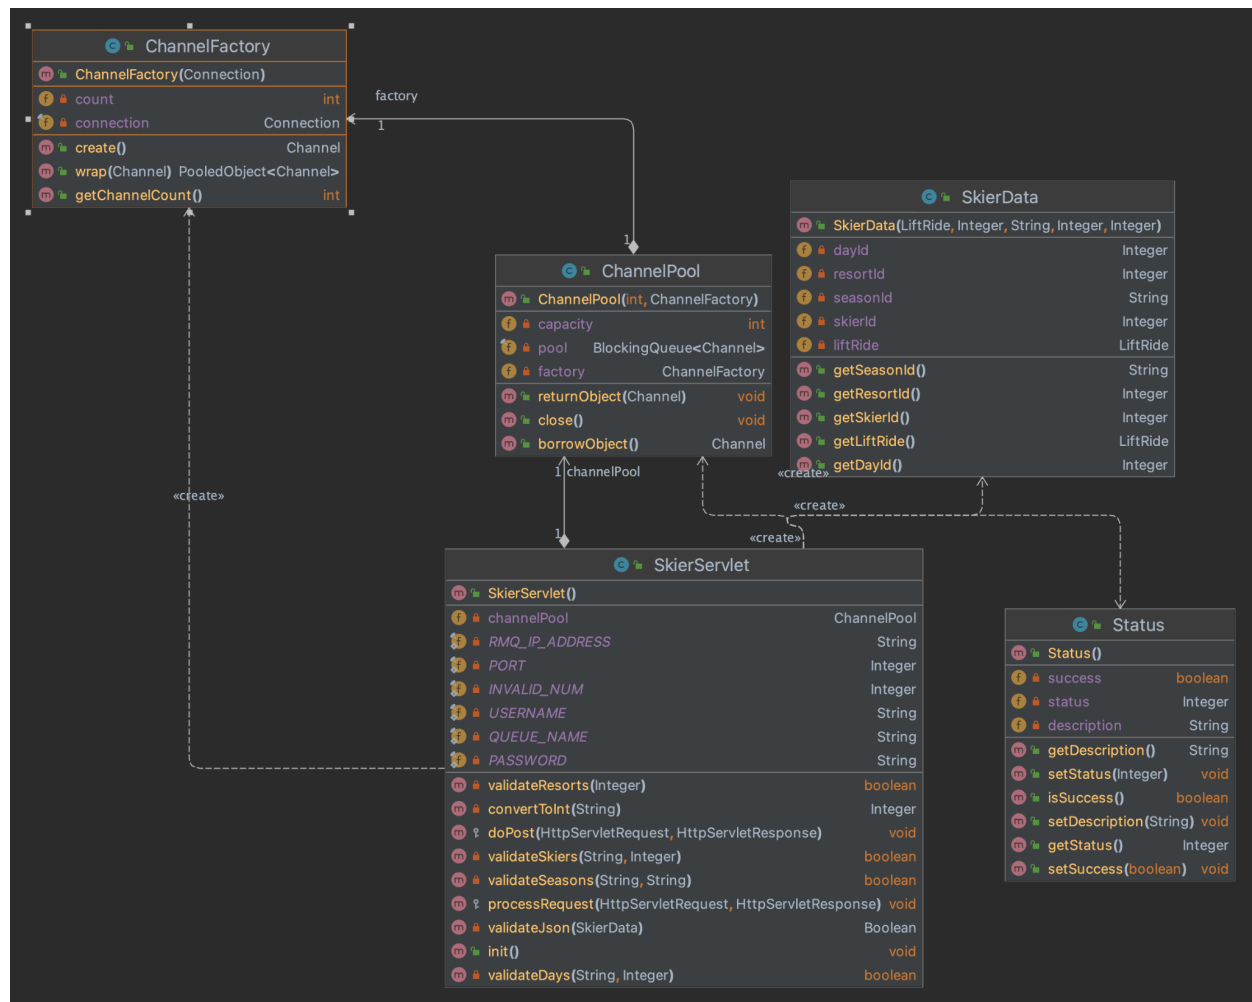


Github URL: <https://github.com/jasongautama/cs6650-Distributed-Systems/tree/main/assignment2>

Class Design

Design for Servlets (Server)

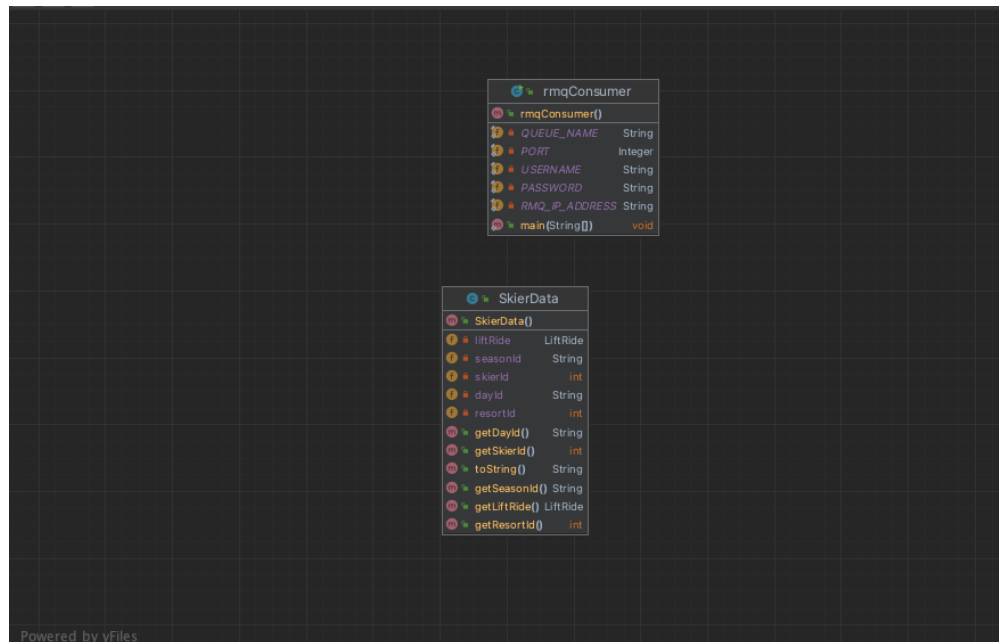
- This class handles the POST request that is validated for the path & LiftRide body when the user sends requests to the EC2 instance.
- Specify the private IP address of the RMQ and keep sending the payload of the **SkierData** to the RMQ
- Created a **ChannelPool.java** so we can reuse the thread whenever they have completed their work.



Design for RabbitMQ

- I installed RabbitMQ in the EC2 instance with IP 35.91.190.94 and then tested the Queue (name *skierQueue*) and was able to send the data to the rabbitMQ
- Success:
 - Setup the Security Group (ports, ip) and access the EC2 instance
 - Installed rabbitMQ to EC2 instance
 - Able to create privilege user to access the management console

Design for Consumer



- The consumer here will retrieve the RabbitMQ *SkierQueue* and store the *SkierData* object into a *ConcurrentHashMap()*
- Create N threads (30, 100, ...) and start consuming the data from the queue

AWS

For EC2 Instance the following is the instances (all using t2.micro):

- 4 x EC2 servlets - accept POST request from client, which is attached to a Load Balancer.
- 1 x EC2 rabbit MQ - queue to store *LiftRide* object from servlets and consumed by consumer
- 1 x EC2 consumer - run `.jar consumer java` file that consume from rabbit MQ

EC2 Global View

Events

Tags

Limits

Instances (6) Info

Instance state = running

Clear filters

Refresh

Connect

Instance state ▾

Actions ▾

Launch Instances

1

Instances

Instances New

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances New

Dedicated Hosts

Scheduled Instances

Capacity Reservations

<input type="checkbox"/>	Name ▲	Instance ID	Instance state ▾	Instance type ▾	Status check	Alarm status	Availability Zone ▾
<input type="checkbox"/>	a2-consumer	i-080a4a4d5d05ae6d5	Running	t2.micro	2/2 checks passed	No alarms +	us-west-2a
<input type="checkbox"/>	a2-rabbitmq-ubuntu	i-0424ddef119862ae6	Running	t2.micro	2/2 checks passed	No alarms +	us-west-2a
<input type="checkbox"/>	a2-server-servlets-DONT-DELETE	i-0bdc09645dbf965ca	Running	t2.micro	2/2 checks passed	No alarms +	us-west-2a
<input type="checkbox"/>	servlets-1	i-072ae014755e26c37	Running	t2.micro	2/2 checks passed	No alarms +	us-west-2c
<input type="checkbox"/>	servlets-2	i-05403fa94c4c383f4	Running	t2.micro	2/2 checks passed	No alarms +	us-west-2c
<input type="checkbox"/>	servlets-3	i-0b309f6abb6890c55	Running	t2.micro	2/2 checks passed	No alarms +	us-west-2c

[Test Case Scenario] No Load Balancer

Explanation: Looking at the RabbitMQ management console, we can see how the message rates plateau around 3000 msg/s. The number indicates the throughput of the message getting received by the queue. Also, the Publish/Consumer-ack have more or less equal number, which is what we want

Phase 1: 32 threads @ 1,000 request / thread

Phase 2: 100 threads @ 840 request / thread

Phase 3: 100 threads @ 840 request / thread

of Consumer thread: 30 threads

==== CLIENT PART 2 ====

mean response time: 29ms

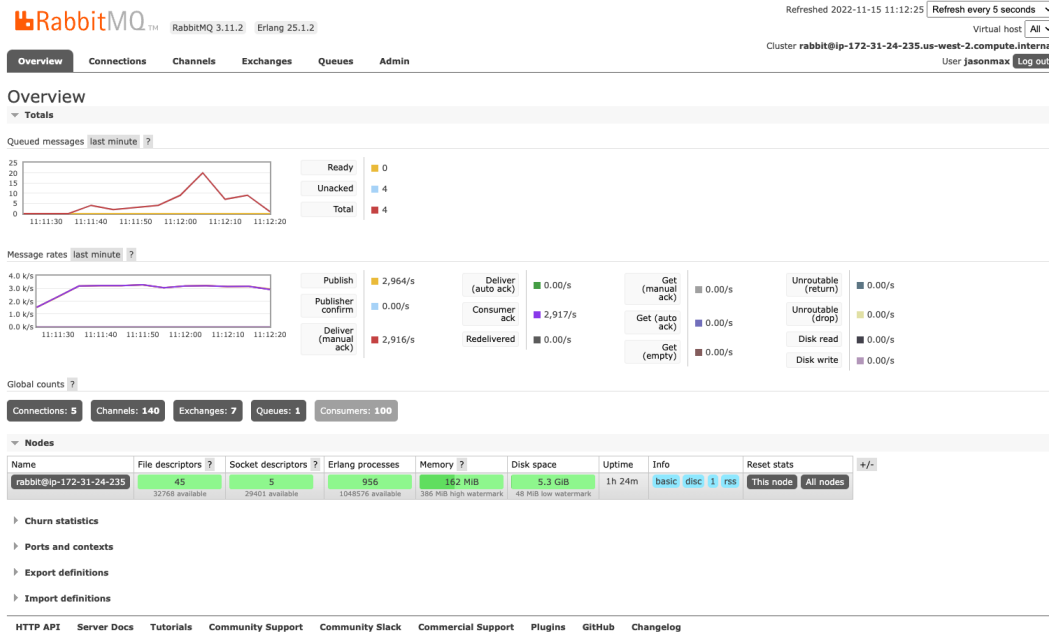
median response time: 29ms

throughput: 2,777req/s

p99 (99th percentile): 40ms

min response time: 12ms

max response time: 219ms



```
Run: part2.Main
/Library/Java/JavaVirtualMachines/jdk-17.0.4.jdk/Contents/Home/bin/java ...

### Start of Phase 1 ###
Produce part1.SkierData to Buffer...
ThreadPool created and executing with numOfThread = 32 & requestPerThread = 1000
...Completed adding part1.SkierData to Buffer
### End of Phase 1 ###

### Start of Phase 2 ###
ThreadPool created and executing with numOfThread = 100 & requestPerThread = 840
### End of Phase 2 ###

### Start of Phase 3 ###
ThreadPool created and executing with numOfThread = 100 & requestPerThread = 840
### End of Phase 3 ###

==== CLIENT PART 2 ====
mean response time: 29ms
median response time: 29ms
throughput: 2777req/s
p99 (99th percentile): 40ms
min response time: 12ms
max response time: 219ms

Process finished with exit code 0
```

[Test Case Scenario] Load Balancer (4 x EC2 t2.micro Instances)

Test with Load Balancer - t2.micro instances x4

Explanation: with 3 additional servlets, we can see that the **throughput increased by 30%** in comparison to a single servlet. From the RMQ console, we can see that the message rates increase to plateau at 4.5k/s (*50% increase in message rates*) as well as the publish/consumer ack is 1.5x faster.

In order to meet the baseline where $\text{production_rate} \approx \text{consumption_rate}$, I increased the amount of thread from 30 to 100 threads on the consumer side to retrieve from the queue to reduce the amount of messages waiting in the Queue. If there's not enough threads to consume the queue, we will see a spike in the graph (ie \wedge), which tells me it was not consuming in equal amount it's producing.

Phase 1: 32 threads @ 1,000 request / thread
Phase 2: 100 threads @ 840 request / thread
Phase 3: 100 threads @ 840 request / thread

of Consumer thread: 100 threads

==== CLIENT PART 2 ====

mean response time: 20ms

median response time: 20ms

throughput: 3636req/s

p99 (99th percentile): 27ms

min response time: 13ms

max response time: 237ms

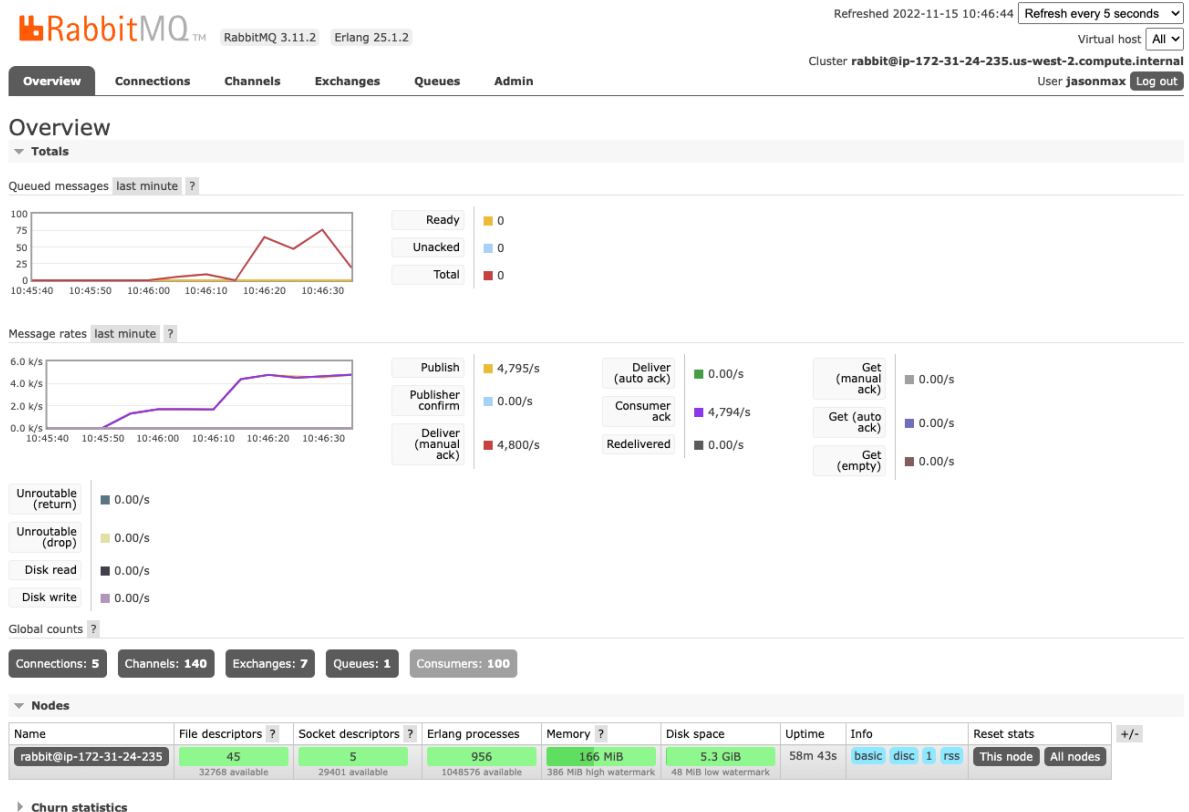


Image showing the RMQ management console where the rate of producing/consuming are almost equal

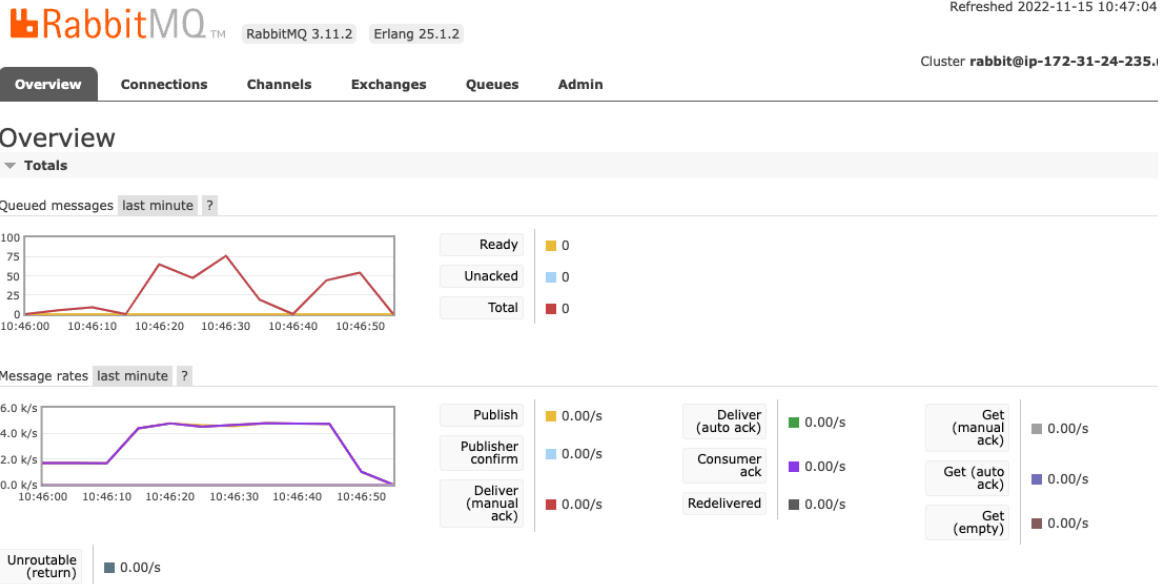
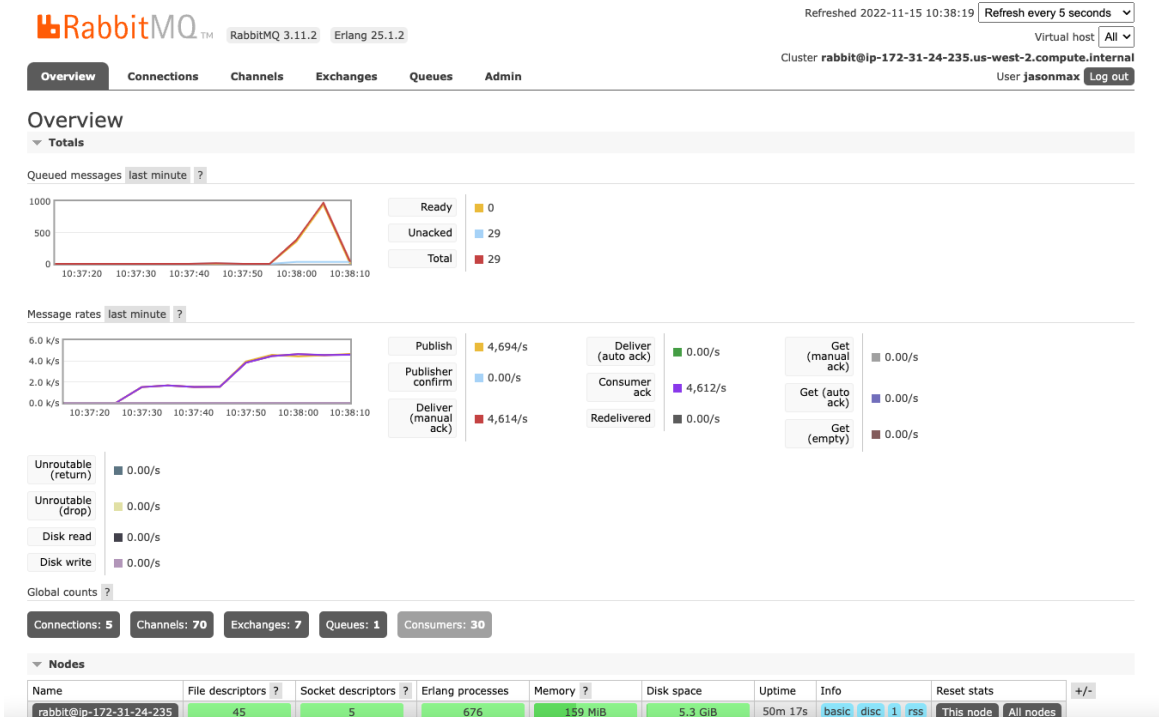


Image showing the RMQ management console receiving messages from 4 t2.micro EC2 instances with the use of Load Balancer



Example where not enough consumer consume the queue that cause a spike ^ in the RMQ